

Assignment2.1

September 16, 2025

1 Assignment 2.1 Use Case - Tayko Software Cataloger

```
[365]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import dmba

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression, LogisticRegressionCV
from sklearn.metrics import accuracy_score, roc_auc_score, confusion_matrix, \
    classification_report
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.tree import DecisionTreeRegressor
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_absolute_error
from dmba import backward_elimination, forward_selection, stepwise_selection

from dmba import regressionSummary
from dmba import adjusted_r2_score, AIC_score, BIC_score
from dmba import classificationSummary
%matplotlib inline
```

Before we start to build predictive models, we first acknowledge that Tayko, a software catalog firm that sells games and educational software, has already conducted a mailing experiment and collected responses from 2,000 customers. This case study examines how the company prepared to launch a revised collection of items in a new catalog mailing and leverages the experimental dataset to model purchasing behavior and spending. The goal is to use these insights to guide Tayko's future catalog campaigns toward higher efficiency and profitability.

```
[366]: df=pd.read_csv('Tayko.csv')

# Basic information
print("Shape of dataset:", df.shape)
print("\nData types:\n", df.dtypes)
df.head()
```

Shape of dataset: (2000, 25)

Data types:

```
sequence_number    int64
US                 int64
source_a           int64
source_c           int64
source_b           int64
source_d           int64
source_e           int64
source_m           int64
source_o           int64
source_h           int64
source_r           int64
source_s           int64
source_t           int64
source_u           int64
source_p           int64
source_x           int64
source_w           int64
Freq               int64
last_update_days_ago int64
1st_update_days_ago int64
Web order          int64
Gender=male        int64
Address_is_res     int64
Purchase           int64
Spending           int64
dtype: object
```

```
[366]:  sequence_number  US  source_a  source_c  source_b  source_d  source_e  \
0              1    1         0         0         1         0         0
1              2    1         0         0         0         0         1
2              3    1         0         0         0         0         0
3              4    1         0         1         0         0         0
4              5    1         0         1         0         0         0

      source_m  source_o  source_h  ...  source_x  source_w  Freq  \
0           0         0         0  ...         0         0     2
1           0         0         0  ...         0         0     0
2           0         0         0  ...         0         0     2
3           0         0         0  ...         0         0     1
4           0         0         0  ...         0         0     1

      last_update_days_ago  1st_update_days_ago  Web order  Gender=male  \
0                3662                3662         1         0
1                2900                2900         1         1
```

2	3883	3914	0	0
3	829	829	0	1
4	869	869	0	0

	Address_is_res	Purchase	Spending
0	1	1	128
1	0	0	0
2	0	1	127
3	0	0	0
4	0	0	0

[5 rows x 25 columns]

```
[367]: df.describe()
```

```
[367]:
```

	sequence_number	US	source_a	source_c	source_b	\
count	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	
mean	1000.500000	0.824500	0.126500	0.056000	0.060000	
std	577.494589	0.380489	0.332495	0.229979	0.237546	
min	1.000000	0.000000	0.000000	0.000000	0.000000	
25%	500.750000	1.000000	0.000000	0.000000	0.000000	
50%	1000.500000	1.000000	0.000000	0.000000	0.000000	
75%	1500.250000	1.000000	0.000000	0.000000	0.000000	
max	2000.000000	1.000000	1.000000	1.000000	1.000000	

	source_d	source_e	source_m	source_o	source_h	...	\
count	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	...	
mean	0.041500	0.151000	0.01650	0.033500	0.052500	...	
std	0.199493	0.358138	0.12742	0.179983	0.223089	...	
min	0.000000	0.000000	0.00000	0.000000	0.000000	...	
25%	0.000000	0.000000	0.00000	0.000000	0.000000	...	
50%	0.000000	0.000000	0.00000	0.000000	0.000000	...	
75%	0.000000	0.000000	0.00000	0.000000	0.000000	...	
max	1.000000	1.000000	1.00000	1.000000	1.000000	...	

	source_x	source_w	Freq	last_update_days_ago	\
count	2000.000000	2000.000000	2000.000000	2000.000000	
mean	0.018000	0.137500	1.417000	2155.101000	
std	0.132984	0.344461	1.405738	1141.302846	
min	0.000000	0.000000	0.000000	1.000000	
25%	0.000000	0.000000	1.000000	1133.000000	
50%	0.000000	0.000000	1.000000	2280.000000	
75%	0.000000	0.000000	2.000000	3139.250000	
max	1.000000	1.000000	15.000000	4188.000000	

	1st_update_days_ago	Web order	Gender=male	Address_is_res	\
count	2000.000000	2000.000000	2000.000000	2000.000000	

mean	2435.601500	0.426000	0.524500	0.221000
std	1077.872233	0.494617	0.499524	0.415024
min	1.000000	0.000000	0.000000	0.000000
25%	1671.250000	0.000000	0.000000	0.000000
50%	2721.000000	0.000000	1.000000	0.000000
75%	3353.000000	1.000000	1.000000	0.000000
max	4188.000000	1.000000	1.000000	1.000000

	Purchase	Spending
count	2000.000000	2000.000000
mean	0.500000	102.62500
std	0.500125	186.78261
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.500000	2.000000
75%	1.000000	153.00000
max	1.000000	1500.00000

[8 rows x 25 columns]

```
[368]: display(df.columns)
```

```
Index(['sequence_number', 'US', 'source_a', 'source_c', 'source_b', 'source_d',
      'source_e', 'source_m', 'source_o', 'source_h', 'source_r', 'source_s',
      'source_t', 'source_u', 'source_p', 'source_x', 'source_w', 'Freq',
      'last_update_days_ago', '1st_update_days_ago', 'Web order',
      'Gender=male', 'Address_is_res', 'Purchase', 'Spending'],
      dtype='object')
```

```
[369]: df.isnull().sum() # Check for missing values
```

```
[369]: sequence_number    0
      US                  0
      source_a            0
      source_c            0
      source_b            0
      source_d            0
      source_e            0
      source_m            0
      source_o            0
      source_h            0
      source_r            0
      source_s            0
      source_t            0
      source_u            0
      source_p            0
      source_x            0
      source_w            0
```

```

Freq                0
last_update_days_ago  0
1st_update_days_ago  0
Web_order           0
Gender=male         0
Address_is_res       0
Purchase            0
Spending            0
dtype: int64

```

1.1 1. Gross Profit = Expect Revenue - Mailing Cost

```

[370]: # 1. Computer the average spending per person in the test mailing
# (Including purchase and non purchased, non purchaser spent = 0)
ave_spending = df['Spending'].mean()

# 2. Get the total revenue which is the answer multiply 180,000

expected_revenue = ave_spending * 180_000
mailing_cost = 180_000 * 2

# 3. Subtract the mailing cost: 180,000 x 2 = $360,000
gross_profit = expected_revenue - mailing_cost

print('Average spending per person:', ave_spending)
print('Expected revenue:', expected_revenue)
print('mailing cost:', mailing_cost)
print('Gross profit:', gross_profit)

```

```

Average spending per person: 102.625
Expected revenue: 18472500.0
mailing cost: 360000
Gross profit: 18112500.0

```

```

[371]: df["US"].value_counts()

```

```

[371]: US
1      1649
0       351
Name: count, dtype: int64

```

```

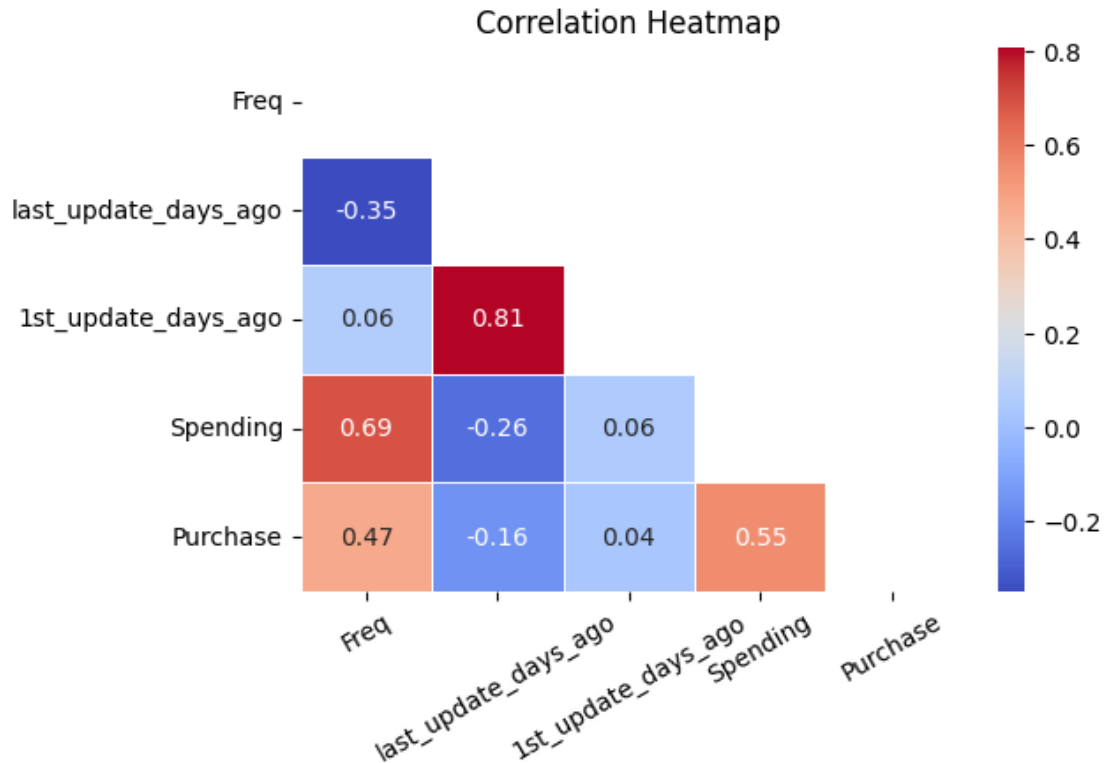
[372]: numeric_features = ["Freq", "last_update_days_ago", "1st_update_days_ago",
↪ "Spending", "Purchase"]
df_numeric = df[numeric_features]
corr = df_numeric.corr()

# Upper Triangle plot

```

```
mask = np.triu(np.ones_like(corr, dtype=bool))

plt.figure(figsize=(6,4))
sns.heatmap(corr, mask=mask, cmap="coolwarm", annot=True,
            annot_kws={"size": 10}, fmt=".2f", linewidths=0.5)
plt.xticks(rotation=30)
plt.title("Correlation Heatmap")
plt.show()
```



1.2 2. Logistic regression Modeling (Model for classifying a customer as a purchaser or nonpurchaser)

1.2.1 2.1 Train/Validation/Test Split (Stratified 800 / 700 / 500)

```
[373]: # drop features
X = df.drop(columns=["Spending", "Purchase", "sequence_number"],
            errors="ignore")
y = df["Purchase"]
# train/valid split
train_X, tmp_X, train_y, tmp_y = train_test_split(X, y, train_size=800,
            stratify=y, random_state=42)
# valid/test split
```

```

valid_X, test_X, valid_y, test_y = train_test_split(tmp_X, tmp_y, test_size=500,
↳stratify=tmp_y, random_state=42)

print('train_X shape:', train_X.shape[0])
print('valid_X shape:', valid_X.shape[0])
print('test_X shape:', test_X.shape[0])

```

```

train_X shape: 800
valid_X shape: 700
test_X shape: 500

```

1.2.2 2.2 Logistic Regression (L2, lbfgs) — Trained on Training Set; Probabilities & Validation Performance

```

[374]: from dmba import classificationSummary
# Build pipeline: standardize features, then logistic regression
pipe = Pipeline([("scaler", StandardScaler()),
                  ("logit", LogisticRegression(penalty="l2", max_iter=1000,
↳solver="lbfgs"))])

# Train model (training set only)
pipe.fit(train_X, train_y)

# probabilities & predictions
proba = pipe.predict_proba(valid_X)          # shape (n, 2) -> columns for
↳classes [0, 1]
pred = pipe.predict(valid_X)

# acc = accuracy_score(valid_y, val_pred)
# auc = roc_auc_score(valid_y, val_proba)

# build result frame
full_result = pd.DataFrame({
    'actual':    valid_y.values,
    'p(0)':      proba[:, 0],
    'p(1)':      proba[:, 1],    # purchase probability
    'predicted': pred
}).sort_values('p(1)', ascending=False)

classificationSummary(full_result.actual, full_result.predicted)
print("classification_report:\n")
print(classification_report(valid_y, pred, digits=3))

```

Confusion Matrix (Accuracy 0.8143)

	Prediction	
Actual	0	1
0	283	67

```

1 63 287
classification_report:

              precision    recall  f1-score   support

0           0.818       0.809       0.813       350
1           0.811       0.820       0.815       350

 accuracy              0.814       700
 macro avg           0.814       0.814       0.814       700
weighted avg           0.814       0.814       0.814       700

```

Interpretation 2.1 and 2.2: the model achieves about 81.4% accuracy, with precision=0.811 and recall=0.820 for purchasers, which means it correctly identifies most buyers while keeping false positives relatively low. The confusion matrix shows that 67 non-purchasers would be mistakenly targeted (wasted mailings), while 63 true purchasers would be missed. The model does a good job ranking which customers are most likely to buy. It performs much better than sending catalogs at random and provides purchase probabilities that can be used to guide profit analysis.

1.3 3. Models for predicting spending among the purchasers

1.3.1 3.1 Purchaser only subsets

From previous defined training and validation sets, keep only records where Purchase = 1. Define predictors X and target y = Spending.

```
[375]: print(train_y.value_counts())
```

```

Purchase
1      400
0      400
Name: count, dtype: int64

```

```
[376]: # 3.1: Purchaser-only subsets; redefine X and y for Spending

# filters for train/validation rows
train_filter = df.index.isin(train_X.index)
valid_filter = df.index.isin(valid_X.index)

# Keep only purchasers (Purchase == 1)
train_purch = df[train_filter & (df["Purchase"] == 1)]
valid_purch = df[valid_filter & (df["Purchase"] == 1)]

# Define predictors (drop targets/ID-like columns) and new target = Spending
X_train_p = train_purch.drop(columns=["Purchase", "Spending", "
↪sequence_number"], errors="ignore")
y_train_p = train_purch["Spending"]

```



```

X_valid_p = valid_purch.drop(columns=["Purchase", "Spending", "sequence_number"], errors="ignore")
y_valid_p = valid_purch["Spending"]

# Quick check
print("Train purchasers:", X_train_p.shape, y_train_p.shape)
print("Valid purchasers:", X_valid_p.shape, y_valid_p.shape)

```

```

Train purchasers: (400, 22) (400,)
Valid purchasers: (350, 22) (350,)

```

```
[377]: y_train_p
```

```

[377]: 2      127
      14      192
      23      174
      29      354
      45      159
      ...
      1972      320
      1980       98
      1986      145
      1989     1030
      1994      184
      Name: Spending, Length: 400, dtype: int64

```

3.2.1 Multiple linear regression

```

[378]: lr = LinearRegression()
lr.fit(X_train_p, y_train_p)
# print coefficients
print('intercept', lr.intercept_)
coef_table = pd.DataFrame({"Predictor": X_train_p.columns,
                           "Coefficient": lr.coef_})
display(coef_table)
# Print performance measures
# # based on the training set (purchasers)
regressionSummary(y_train_p, lr.predict(X_train_p))

```

```
intercept 122.90673571037372
```

	Predictor	Coefficient
0	US	18.618756
1	source_a	22.970667
2	source_c	-43.699780
3	source_b	-37.639370
4	source_d	-65.071902
5	source_e	-39.933289
6	source_m	-55.665131

7	source_o	41.312097
8	source_h	-152.976257
9	source_r	49.381204
10	source_s	-13.427969
11	source_t	-64.363563
12	source_u	13.002463
13	source_p	-59.026222
14	source_x	-33.834799
15	source_w	4.796471
16	Freq	84.594789
17	last_update_days_ago	-0.022070
18	1st_update_days_ago	-0.002507
19	Web order	-3.931550
20	Gender=male	-35.643459
21	Address_is_res	-90.347336

Regression statistics

Mean Error (ME) : 0.0000
Root Mean Squared Error (RMSE) : 181.8183
Mean Absolute Error (MAE) : 108.5417
Mean Percentage Error (MPE) : -83.3628
Mean Absolute Percentage Error (MAPE) : 112.4249

```
[379]: # Get predictions based on X_valid_p
lin_pred = lr.predict(X_valid_p)

# Get RMSE and MAE on validation set
lin_rmse = np.sqrt(mean_squared_error(y_valid_p, lin_pred))
lin_mae = mean_absolute_error(y_valid_p, lin_pred)

result = pd.DataFrame({'Predicted': lin_pred, 'Actual': y_valid_p,
                      'Residual': y_valid_p - lin_pred})

print(result.head(20))

# based on the validation set
regressionSummary(y_valid_p, lin_pred)
```

	Predicted	Actual	Residual
0	88.797490	128	39.202510
8	444.408741	489	44.591259
9	170.077343	174	3.922657
18	94.592360	130	35.407640
20	283.733158	386	102.266842
21	276.931351	161	-115.931351
24	229.195120	131	-98.195120
25	147.620586	189	41.379414

31	175.242849	352	176.757151
41	207.158200	34	-173.158200
43	830.957907	639	-191.957907
44	253.905185	638	384.094815
50	241.647766	232	-9.647766
54	170.793172	375	204.206828
60	244.617923	136	-108.617923
78	201.305580	161	-40.305580
81	168.674359	129	-39.674359
83	238.870108	98	-140.870108
86	199.759076	158	-41.759076
105	147.586064	405	257.413936

Regression statistics

Mean Error (ME) : -21.1440
 Root Mean Squared Error (RMSE) : 144.9845
 Mean Absolute Error (MAE) : 99.8359
 Mean Percentage Error (MPE) : -92.6650
 Mean Absolute Percentage Error (MAPE) : 116.2232

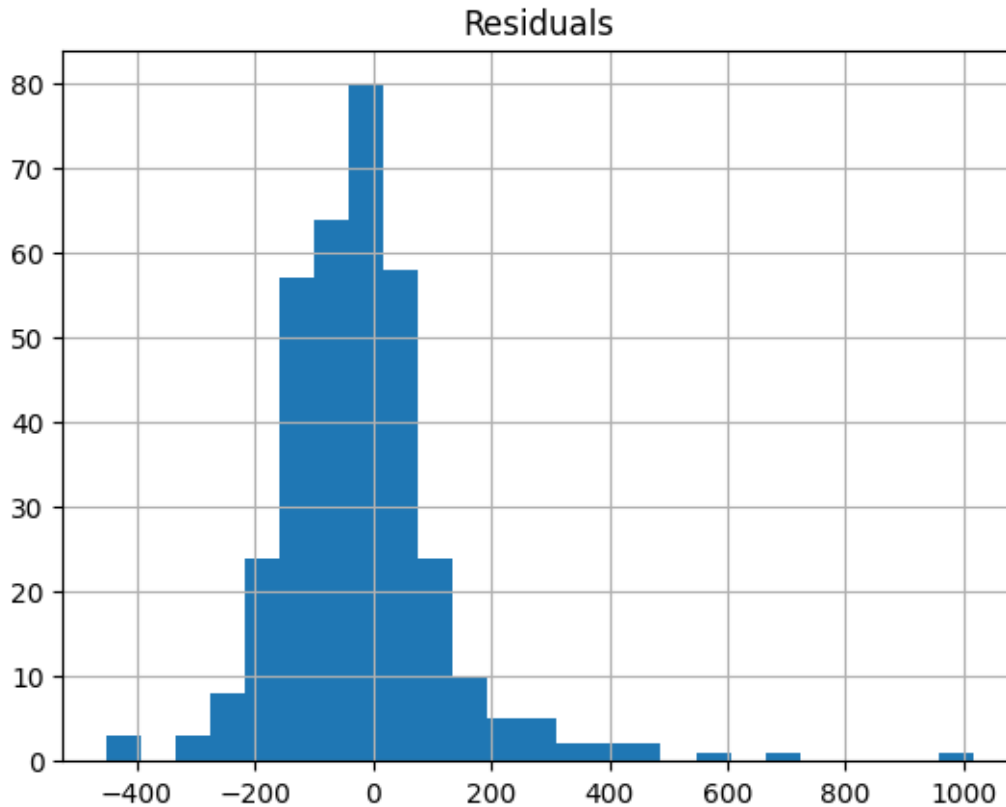
```
[380]: # Determine the percentage of datapoints with a residual in 75%
        ↪ threshold
residuals = y_valid_p - lin_pred
# Take absolute value
abs_resid = np.abs(residuals)

# 75th percentile threshold
threshold = np.percentile(abs_resid, 75)

# Percentage of data points within this threshold
within_75 = np.mean(abs_resid <= threshold) * 100

print("75% residual threshold:", threshold)
print(len(residuals[(residuals > -134) & (residuals < 134)]) / len(residuals))
pd.DataFrame({"Residuals": residuals}).hist(bins=25)
plt.show()
```

75% residual threshold: 134.40496491556274
 0.7457142857142857



Most prediction errors are centered around zero with a roughly symmetric distribution, and about 75% of the residuals fall within a moderate range, means the model good fit with some larger outliers.

3.2.2 Regression trees

```
[ ]: tree = DecisionTreeRegressor(max_depth=5, random_state=42)
tree.fit(X_train_p, y_train_p)

tree_pred = tree.predict(X_valid_p)
tree_rmse = np.sqrt(mean_squared_error(y_valid_p, tree_pred))
tree_mae = mean_absolute_error(y_valid_p, tree_pred)
```

3.2.3 Model selection and explain Compare both models on the validation set, choose the better one.

```
[393]: print("Validation Results:")
print(f"Linear Regression: RMSE={lin_rmse:.2f}, MAE={lin_mae:.2f}")
print(f"Regression Tree: RMSE={tree_rmse:.2f}, MAE={tree_mae:.2f}")

# Choose model with lower error
if tree_rmse < lin_rmse:
```

```

    print("→ Regression Tree performs better on validation set.")
else:
    print("→ Linear Regression performs better on validation set.")

```

Validation Results:

Linear Regression: RMSE=144.98, MAE=99.84

Regression Tree: RMSE=175.23, MAE=104.98

→ Linear Regression performs better on validation set.

stepwise forward selection for the linear regression

```

[383]: from sklearn.feature_selection import SequentialFeatureSelector
       from sklearn.metrics import mean_squared_error, mean_absolute_error

       # Forward stepwise on TRAIN only
       base_lr = LinearRegression()
       sfs = SequentialFeatureSelector(
           base_lr, direction="forward", n_features_to_select="auto",
           scoring="neg_mean_squared_error", cv=5, n_jobs=-1
       ).fit(X_train_p, y_train_p)

       # Selected features and final fit
       sel_feats = X_train_p.columns[sfs.get_support()]
       best_model = LinearRegression().fit(X_train_p[sel_feats], y_train_p)

       # Validate
       lr_pred = best_model.predict(X_valid_p[sel_feats])
       lr_rmse = np.sqrt(mean_squared_error(y_valid_p, lr_pred))
       lr_mae = mean_absolute_error(y_valid_p, lr_pred)

       print("Stepwise LR selected:", list(sel_feats))
       print(f"Validation - Linear Regression: RMSE={lr_rmse:.2f}, MAE={lr_mae:.2f}")

```

Stepwise LR selected: ['source_a', 'source_c', 'source_d', 'source_m',
'source_h', 'source_t', 'source_p', 'Freq', 'last_update_days_ago',
'Gender=male', 'Address_is_res']

Validation - Linear Regression: RMSE=144.79, MAE=98.02

Interpretation 3.2.3: The goal in Question 3 was to predict how much a customer will spend if they purchase. Between the two models tested, multiple linear regression achieved lower errors (RMSE = 144.98, MAE = 99.84) than the regression tree (RMSE = 175.23, MAE = 104.98). Stepwise regression further confirmed the strength of the linear model, and validation results showed it predicts spending more accurately. Therefore, linear regression is chosen as the better model for forecasting spending.

1.4 4. Score Analysis and Profit Estimation Using Test Data

1.4.1 4.1 Add a column - Logistic regression

```
[384]: # Predict probability of purchase (class=1) from logistic regression pipeline
test_proba = pipe.predict_proba(test_X)[: , 1]

# Create Score Analysis DataFrame
score_analysis = test_X.copy()
score_analysis["Purchase"] = test_y.values
score_analysis["p_purchase"] = test_proba
```

```
[385]: score_analysis.dtypes
```

```
[385]: US                                int64
source_a                               int64
source_c                               int64
source_b                               int64
source_d                               int64
source_e                               int64
source_m                               int64
source_o                               int64
source_h                               int64
source_r                               int64
source_s                               int64
source_t                               int64
source_u                               int64
source_p                               int64
source_x                               int64
source_w                               int64
Freq                                   int64
last_update_days_ago                  int64
1st_update_days_ago                   int64
Web order                             int64
Gender=male                           int64
Address_is_res                        int64
Purchase                              int64
p_purchase                            float64
dtype: object
```

1.4.2 4.2 Add another column - predict spending amount from (chosen model from 3.2.3)

```
[386]: # chosen linear regression (stepwise)
# Use the same selected predictors as in stepwise
spend_pred = best_model.predict(test_X[sel_feats])

score_analysis["pred_spending"] = spend_pred
```

```
[387]: list(score_analysis.columns)
score_analysis.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 500 entries, 1052 to 475
Data columns (total 25 columns):
#   Column                Non-Null Count  Dtype
---  -
0   US                     500 non-null    int64
1   source_a               500 non-null    int64
2   source_c               500 non-null    int64
3   source_b               500 non-null    int64
4   source_d               500 non-null    int64
5   source_e               500 non-null    int64
6   source_m               500 non-null    int64
7   source_o               500 non-null    int64
8   source_h               500 non-null    int64
9   source_r               500 non-null    int64
10  source_s               500 non-null    int64
11  source_t               500 non-null    int64
12  source_u               500 non-null    int64
13  source_p               500 non-null    int64
14  source_x               500 non-null    int64
15  source_w               500 non-null    int64
16  Freq                   500 non-null    int64
17  last_update_days_ago   500 non-null    int64
18  1st_update_days_ago    500 non-null    int64
19  Web order              500 non-null    int64
20  Gender=male            500 non-null    int64
21  Address_is_res         500 non-null    int64
22  Purchase               500 non-null    int64
23  p_purchase             500 non-null    float64
24  pred_spending          500 non-null    float64
dtypes: float64(2), int64(23)
memory usage: 101.6 KB
```

```
[388]: score_analysis.head()
```

```
[388]:
```

	US	source_a	source_c	source_b	source_d	source_e	source_m	\
1052	1	1	0	0	0	0	0	
1603	1	0	0	0	0	0	0	
1784	0	0	0	0	0	0	0	
464	0	0	0	0	0	0	0	
752	1	1	0	0	0	0	0	

	source_o	source_h	source_r	...	source_w	Freq	last_update_days_ago	\
1052	0	0	0	...	0	1	3067	
1603	0	0	0	...	0	1	2690	

1784	0	0	0 ...	0	1	4127
464	0	0	0 ...	1	1	1091
752	0	0	0 ...	0	1	2947

	1st_update_days_ago	Web order	Gender=male	Address_is_res	Purchase	\
1052	3067	0	0	0	1	
1603	2690	1	0	0	1	
1784	4127	0	0	0	1	
464	1091	0	0	0	0	
752	2947	0	1	0	1	

	p_purchase	pred_spending
1052	0.581226	167.637028
1603	0.770639	148.687872
1784	0.551669	113.068412
464	0.409325	188.322887
752	0.496562	140.151103

[5 rows x 25 columns]

1.4.3 4.3 skip

1.4.4 4.4 Add expected spending (adjusted probability \times predicted spending)

1. In the mailing experiment, Tayko mailed 20,000 names and got 1,065 purchasers. That gives the true response rate: $1065 / 20000 = 0.053$
2. But for modeling, the dataset was rebalanced to 1,000 purchasers and 1,000 non-purchasers. That makes the apparent response rate = 0.5 (50%) in the training data.
3. So after fitting the logistic regression, we need to adjust the model's predicted purchase probability back to the true population rate to get the adjusted probability of purchase $0.053/0.5 = 0.107$.

```
[389]: adjust_factor = 0.053 / 0.5    # = 0.107
score_analysis["expected_spending"] = (
    score_analysis["p_purchase"] * adjust_factor *
    ↪ score_analysis["pred_spending"]
)
```

1.4.5 4.5 Plot cumulative gains chart of the expected spending

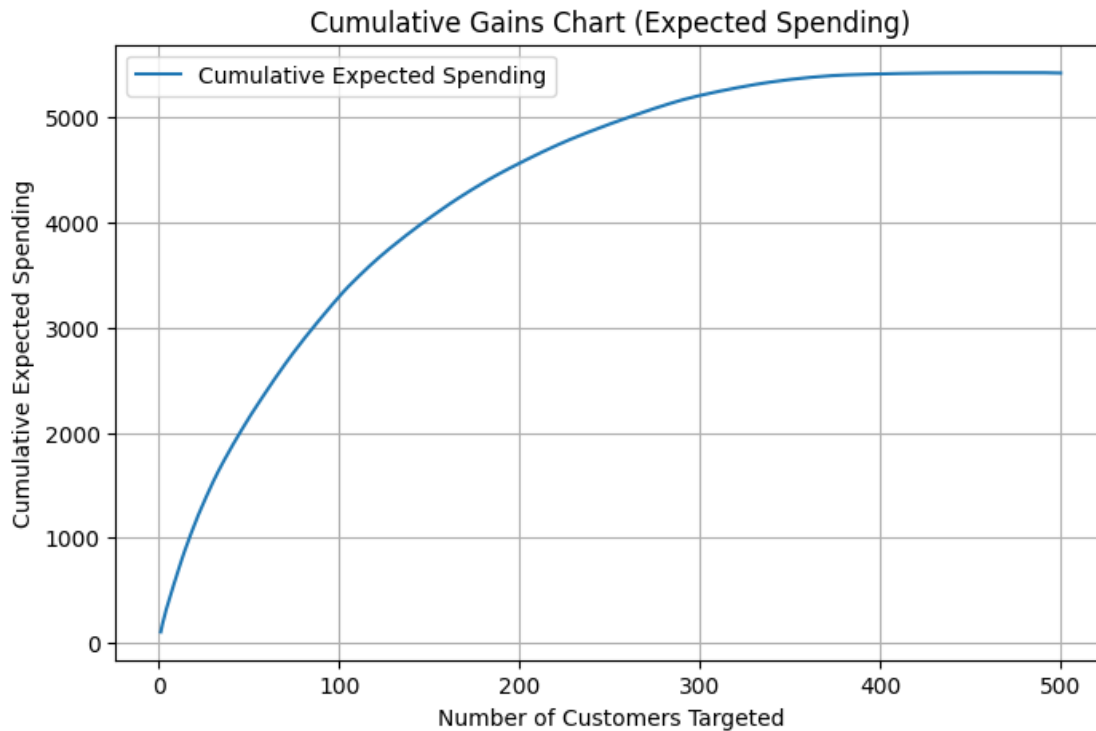
```
[390]: # Sort customers by expected spending (highest to lowest)
score_sorted = score_analysis.sort_values("expected_spending", ascending=False).
    ↪ reset_index(drop=True)

# Cumulative expected spending
cumulative_spending = np.cumsum(score_sorted["expected_spending"])

plt.figure(figsize=(8,5))
```



```
plt.plot(range(1, len(cumulative_spending)+1), cumulative_spending,
        label="Cumulative Expected Spending")
plt.xlabel("Number of Customers Targeted")
plt.ylabel("Cumulative Expected Spending")
plt.title("Cumulative Gains Chart (Expected Spending)")
plt.legend()
plt.grid(True)
plt.show()
```



1.4.6 4.6 using this cumulative gains curve (not directly), estimate the gross profit - based on the model.

Estimate gross profit for mailing 180,000 names

1. Compute the average expected spending per record in the test sample.
2. Scale to 180,000.
3. Subtract mailing costs ($\$2 \times 180,000$).

```
[391]: # Average expected spending in test sample
avg_exp_spending = score_analysis["expected_spending"].mean()

# Scale up to 180,000 customers
total_expected_spending = avg_exp_spending * 180000
```

```
# Mailing cost
mailing_cost = 180000 * 2

# Gross profit
gross_profit = total_expected_spending - mailing_cost

print(f"Avg expected spending per record: ${avg_exp_spending:.2f}")
print(f"Total expected spending: ${total_expected_spending:.2f}")
print(f"Total expected cost: ${mailing_cost:.2f}")
print(f"Estimated gross profit (180k mailing): ${gross_profit:,.2f}")
```

Avg expected spending per record: \$10.84
Total expected spending: \$1951953.83
Total expected cost: \$360000.00
Estimated gross profit (180k mailing): \$1,591,953.83

```
[392]: score_analysis.head()
```

```
[392]:
```

	US	source_a	source_c	source_b	source_d	source_e	source_m	\
1052	1	1	0	0	0	0	0	
1603	1	0	0	0	0	0	0	
1784	0	0	0	0	0	0	0	
464	0	0	0	0	0	0	0	
752	1	1	0	0	0	0	0	

	source_o	source_h	source_r	...	Freq	last_update_days_ago	\
1052	0	0	0	...	1	3067	
1603	0	0	0	...	1	2690	
1784	0	0	0	...	1	4127	
464	0	0	0	...	1	1091	
752	0	0	0	...	1	2947	

	1st_update_days_ago	Web order	Gender=male	Address_is_res	Purchase	\
1052	3067	0	0	0	1	
1603	2690	1	0	0	1	
1784	4127	0	0	0	1	
464	1091	0	0	0	0	
752	2947	0	1	0	1	

	p_purchase	pred_spending	expected_spending
1052	0.581226	167.637028	10.328112
1603	0.770639	148.687872	12.145977
1784	0.551669	113.068412	6.611893
464	0.409325	188.322887	8.171040
752	0.496562	140.151103	7.376933

[5 rows x 26 columns]

1.5 5. Summary and Result

Business Objective As the data scientist supporting Tayko, a software catalog firm that sells games and educational software, my objective is to help optimize its catalog mailing strategy by reducing wasted costs and focusing on customers most likely to purchase and spend. Since sending the catalog to everyone is prohibitively expensive, the goal is to target the right group of customers in order to maximize profit.

Models Used and Why

To do this, I applied a few different data mining approaches. First, I used logistic regression to classify customers as purchasers or non-purchasers and to generate a probability of purchase for each record. Then, for those predicted to purchase, I tested both multiple linear regression and regression trees to estimate spending. I also applied stepwise regression to reduce predictors and avoid overfitting, ensuring that only the most relevant variables were kept in the model. The linear regression model consistently produced lower error rates (RMSE and MAE) compared to the regression tree, and stepwise selection confirmed that it provided a more stable and interpretable solution.

Model Results and Recommendations

The results show that this combined modeling strategy works well. Logistic regression achieved strong classification accuracy, and linear regression predicted spending with better performance than regression trees. The cumulative gains chart further demonstrated that targeting customers by expected spending yields far greater returns than random mailing. Using the predicted probabilities and expected spending, the estimated gross profit from mailing to 180,000 customers is about \$1.59 million after costs.

Recommendation: Based on the gains curve, Tayko should target approximately the top 120,000–140,000 customers ranked by expected spending (probability \times predicted spending), rather than mailing to the entire 180,000 names. This cutoff balances mailing costs against incremental revenue, ensuring each additional mailing remains profitable. Under this strategy, Tayko can expect an estimated \$1.5–\$1.6 million gross profit, a significant improvement over a random or full mailing.

Next Steps: Tayko should operationalize this by (1) applying the logistic regression model to score the entire 200,000 eligible customers, (2) predicting spending with the linear regression model for likely purchasers, (3) ranking all customers by expected spending, and (4) selecting the top-ranked segment (120k–140k) for catalog distribution. This ensures the company reduces wasted costs, focuses on the most profitable customers, and creates a repeatable, data-driven process for future campaigns.